Slide 1: Retrieval Augmented Generation

RAG lets a model look things up before it answers. You embed documents into vectors, store

```
def retrieve(q, k=5):
    return index.search(embed(q), k)
```

Slide 2: Chunking Strategies

If chunks are too big you blow the context window; too small and you lose meaning. Many te

```
def retrieve(q, k=5):
    return index.search(embed(q), k)
```

Slide 3: Evaluation

You cannot improve what you do not measure. Build a golden set of questions and expected a

```
def retrieve(q, k=5):
    return index.search(embed(q), k)
```

Slide 4: The Forward Deployed Engineer Role

An FDE sits with the customer, understands their data and workflow, and ships a working in

```
def retrieve(q, k=5):
    return index.search(embed(q), k)
```

Slide 5: Prompt Caching

Prompt caching reuses a stable prefix so repeated calls are cheaper. Put the system prompt

```
def retrieve(q, k=5):
    return index.search(embed(q), k)
```